


QUAM — Query Affinity Modelling

A learned co-relevance graph guides the next batch to rerank



The graph is **learned, not crawled** — edges encode whether two documents satisfy the same need.

Up to +26% Recall@1000 over the standard cascade · ~3% latency overhead.