


SUNAR — multi-hop QA

F1 score vs the strongest decomposition baselines · LLM-agnostic



*LLM-agnostic — gains hold across GPT-3.5, GPT-4o-mini, Llama, and Mistral.
Uncertainty itself becomes the relevance signal.*