

Why this matters

Comparable task performance, with far fewer expensive calls

HEADLINE



fewer LLM calls

same task performance

LLM CALLS PER QUERY

each ● = 10 expensive calls · same task, same final accuracy

classical baseline



CASE



The same algorithm picks reranker-bound documents, in-context exemplars, evaluation samples — anywhere you choose a subset under a budget.

Better retrieval is useless if it's too expensive.

We're pushing the cost frontier, not just the quality frontier.