

Why this matters

Comparable task performance, with far fewer expensive calls

HEADLINE

7x

fewer LLM calls

same task performance